

Michael Dang – 16257750

COMP-SCI 5565

Final Project

12/16/2023

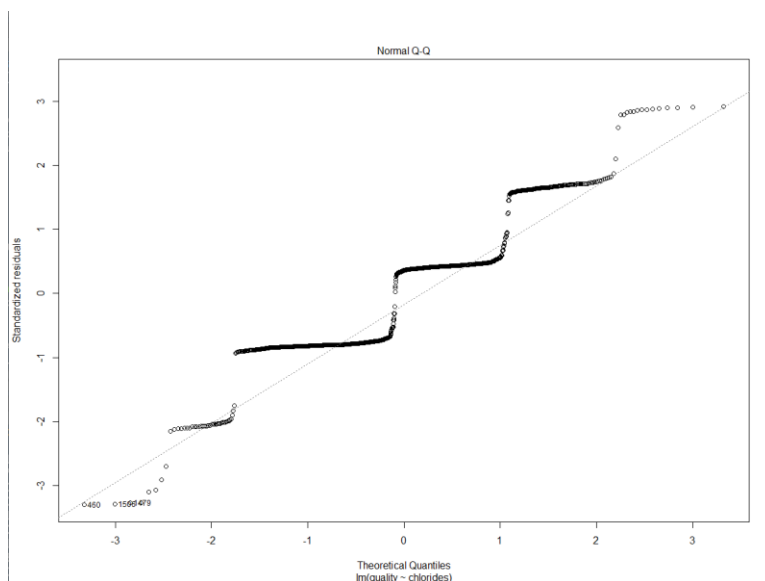
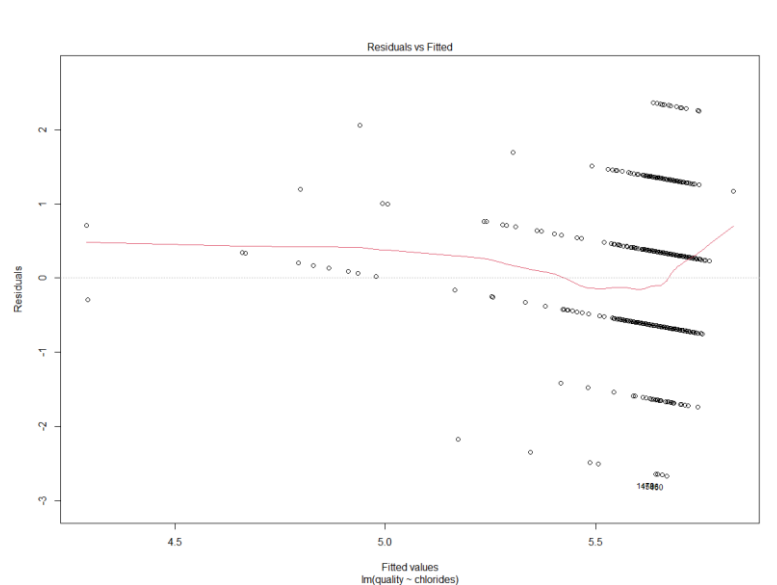
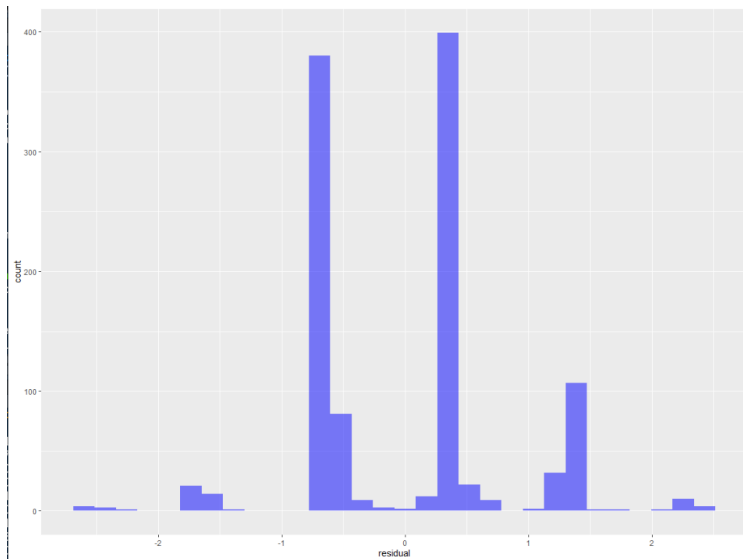
## 1. Regression

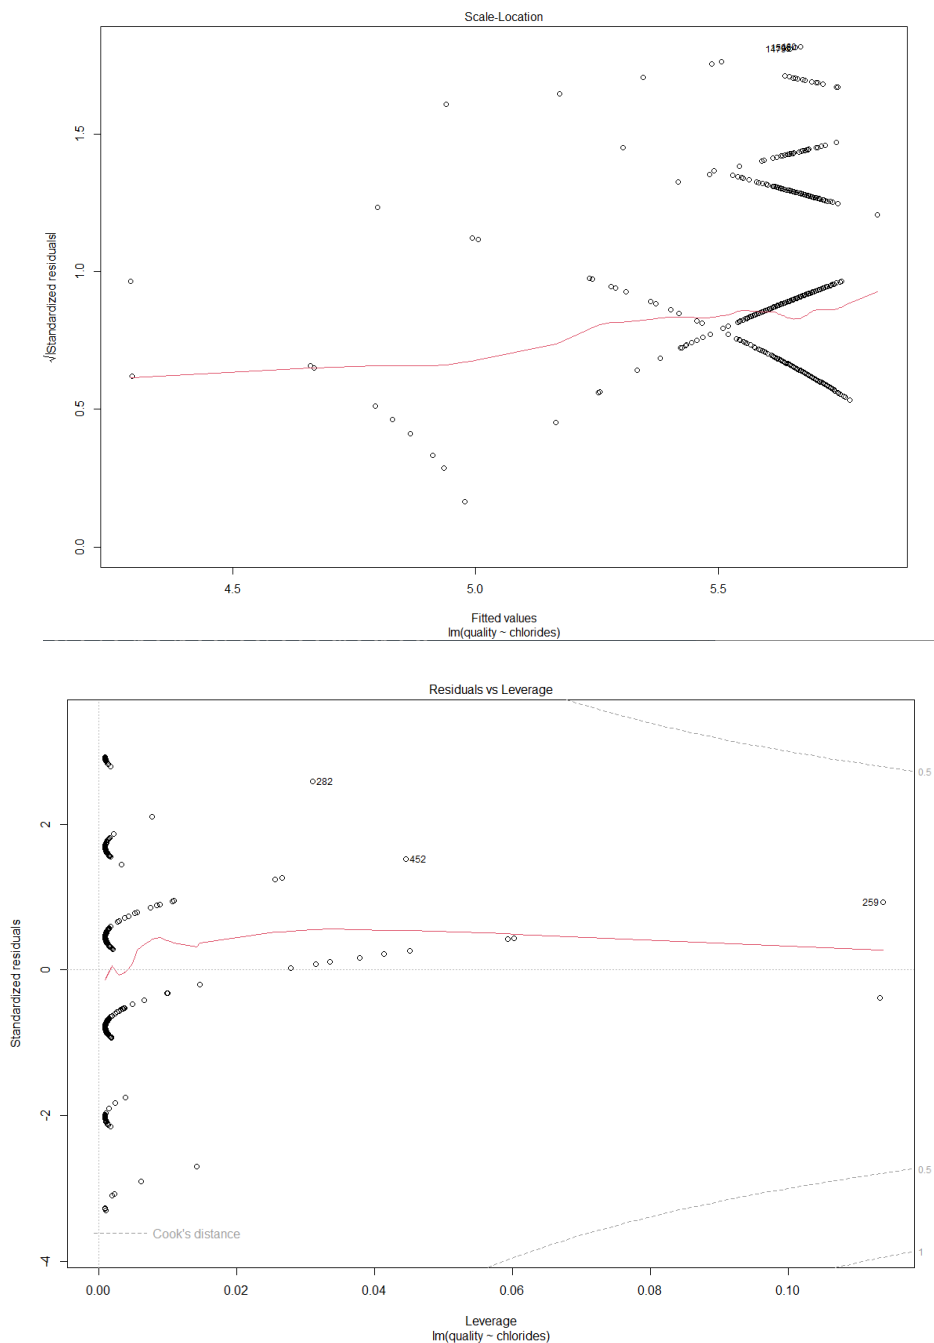
Wine quality dataset

### Linear Regression

```
1 library(caTools)
2 library(ggplot2)
3 #Regression problem
4
5 #Upload the data - Heart disease UCI
6 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\winequality-red.csv")
7
8 #Check N/A value
9 print(any(is.na(data)))
10
11 #Check structure of the data
12 print(str(data))
13
14 #Train validation test split
15 set.seed(42)
16 spec <- c(train = .7, validate = .15, test = .15)
17
18 g <- sample(cut(
19   seq(nrow(data)),
20   nrow(data)*cumsum(c(0,spec)),
21   labels = names(spec)
22 ))
23
24 res <- split(data, g)
25
26 cat('\n Shape of train: \n', dim(res$train), '\n Shape of validation: \n', dim(res$validate), '\n Shape of test: \n', dim(res$test))
27
28 #Linear Regression
29 #Fit the model
30 model <- lm(quality ~ chlorides, data = res$train)
31 print(summary(model))
32
33 #Plot residual check normality
34 #Grab residuals
35 residual <- residuals(model)
36
37 #Convert to dataframe and plot it
38 residual <- as.data.frame(residual)
39 print('Head of the residual :')
40 print(head(residual))
41
42 #ggplot
43 ggplot(residual, aes(residual)) + geom_histogram(fill = 'blue', alpha = 0.5)
44
45 #Predict on the validation set
46 val.prediction <- predict(model, res$validate)
47
48 #Result
49 val.results <- cbind(val.prediction, res$test$quality)
50 colnames(val.results) <- c('pred', 'valid')
51 val.results <- as.data.frame(val.results)
52
53 #Error
54 val_mse <- mean((val.results$valid - val.results$pred)^2)
55 cat('\nValidation MSE: ', val_mse, '\n')
56
57 #R^2 adjusted
58 cat('\nR^2 adjusted ', summary(model)$adj.r.squared, '\n')
59
60 #Predict on the test set
61 cat('\n***** Prediction the test set *****')
62 test_predictions <- predict(model, res$test)
63
64 #Error
65 test_mse <- mean((test_predictions - res$test$quality)^2)
66 cat('\n Test MSE: ', test_mse, '\n')
67
```

We want a histogram of our residuals to be normally distributed, something with a strong bimodal distribution may be a warning that our data was not a good fit for linear regression.





Looks like the data is following a normal distribution.

Code output

```
> source("~/active-rstudio-document")
[1] FALSE
'data.frame': 1599 obs. of 12 variables:
 $ fixed.acidity : num 7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
 $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
 $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
 $ residual.sugar : num 1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 6.1 ...
 $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
 $ free.sulfur.dioxide : num 11 25 15 17 11 13 15 15 9 17 ...
 $ total.sulfur.dioxide : num 34 67 54 60 34 40 59 21 18 102 ...
 $ density : num 0.998 0.997 0.997 0.998 0.998 0.998 ...
 $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
 $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57 0.8 ...
 $ alcohol : num 9.4 9.8 9.8 9.8 9.4 9.4 10 9.5 10.5 ...
 $ quality : int 5 5 5 6 5 5 5 7 5 ...
NULL

Shape of train:
1119 12
Shape of validation:
240 12
Shape of test:
240 12
```

## Model summary

```
Call:
lm(formula = quality ~ chlorides, data = res$train)

Residuals:
    Min       1Q   Median       3Q      Max
-2.6678 -0.6499  0.2912  0.3617  2.3630

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  5.85756    0.05141  113.938  < 2e-16 ***
chlorides   -2.56418    0.51935   -4.937  9.13e-07 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.8094 on 1117 degrees of freedom
Multiple R-squared:  0.02136,    Adjusted R-squared:  0.02048
F-statistic: 24.38 on 1 and 1117 DF,  p-value: 9.126e-07
```

Head of the residual column and MSE for validation and test set.

Notice: I'm using the  $R^2$  adjusted instead of  $R^2$  because it is more accurate

(My statistics professor said so)

```
[1] "Head of the residual :"
```

	residual
1	-0.6626805
2	-0.6062685
4	0.3347553
7	-0.6806298
8	1.3091135
11	-0.6088326

```
Validation MSE:  0.5853875

R^2 adjusted  0.02048154

***** Prediction the test set *****
Test MSE:  0.5568056
> |
```

Interpreting the simple linear regression result shows that the regression model shows good MSE but a bad R-square value. Indicate that the model has a low error when predicting values but there is very little correlation between the variables.

## Polynomial Regression

```
1 library(caTools)
2 library(ggplot2)
3 #Regression problem
4
5 #Upload the data - Heart disease UCI
6 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\winequality-red.csv")
7
8 #Check N/A value
9 print(any(is.na(data)))
10
11 #Check structure of the data
12 print(str(data))
13
14 #Train validation test split
15 set.seed(42)
16 spec <- c(train = .7, validate = .15, test = .15)
17
18 g <- sample(cut(
19   seq(nrow(data)),
20   nrow(data)*cumsum(c(0,spec))),
21   labels = names(spec))
22 ))
```

Data	
• data	1599 obs. of 12 variables
• ploy_model	List of 12
• poly_model	List of 12
• res	List of 3
• residual	1119 obs. of 1 variable
• val.results	240 obs. of 2 variables
Values	
g	Factor w/ 3 levels "train","validate"...
spec	Named num [1:3] 0.7 0.15 0.15
test_mse	0.39387576918608
test_predict...	Named num [1:240] 5.28 5.17 5.4 5.9 5...
val_mse	0.653745606607503
val.predicti...	Named num [1:240] 5.42 5.22 5.23 5.27...

```

24 res <- split(data, g)
25
26 cat('\n Shape of train: \n', dim(res$train), '\n Shape of validation: \n', dim(res$validate), '\n Shape of test: \n', dim(res$test))
27
28 # Polynomial Regression
29 #Fit the model
30 poly_model <- lm(quality ~ poly(chlorides, sulphates, alcohol), data = res$train)
31 print(summary(poly_model))
32
33 #Plot residual check normality
34 #Grab residuals
35 residual <- residuals(poly_model)
36
37 #Convert to dataframe and plot it
38 residual <- as.data.frame(residual)
39 print('Head of the residual :')
40 print(head(residual))
41
42 #ggplot
43 ggplot(residual, aes(residual)) + geom_histogram(fill = 'blue', alpha = 0.5)
44
45 #Predict on the validation set
46 val.prediction <- predict(poly_model, res$validate)
47
48 #Result
49 val.results <- cbind(val.prediction, res$test$quality)
50 colnames(val.results) <- c('pred', 'valid')
51 val.results <- as.data.frame(val.results)
52
53 #Error
54 val_mse <- mean((val.results$valid - val.results$pred)^2)
55 cat('\nValidation MSE: ', val_mse, '\n')
56
57 #RA2 adjusted
58 cat('\nRA2 adjusted ', summary(poly_model)$adj.r.squared, '\n')
59
60 #Predict on the test set
61 cat('\n***** Prediction the test set *****')
62 test_predictions <- predict(poly_model, res$test)
63
64 #Error
65 test_mse <- mean((test_predictions - res$test$quality)^2)
66 cat('\n Test MSE: ', test_mse, '\n')
67 |

```

## Code output

```

R 4.2.1 ~ / .
> source("../.active-rstudio-document")
[1] FALSE
'data.frame': 1599 obs. of 12 variables:
 $ fixed.acidity : num 7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
 $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
 $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
 $ residual.sugar : num 1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 6.1 ...
 $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
 $ free.sulfur.dioxide : num 11 25 15 17 11 13 15 15 9 17 ...
 $ total.sulfur.dioxide : num 34 67 54 60 34 40 59 21 18 102 ...
 $ density : num 0.998 0.997 0.997 0.998 0.998 ...
 $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
 $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57 0.8 ...
 $ alcohol : num 9.4 9.8 9.8 9.8 9.4 9.4 10 9.5 10.5 ...
 $ quality : int 5 5 5 6 5 5 5 7 5 ...
NULL

Shape of train:
1119 12
Shape of validation:
240 12
Shape of test:
240 12

```

## Model summary

```
Call:
lm(formula = quality ~ poly(chlorides, sulphates, alcohol), data = res$train)

Residuals:
    Min       1Q   Median       3Q      Max
-2.64857 -0.36264 -0.09975  0.48912  2.33734

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      5.63360    0.02074   271.65 < 2e-16 ***
poly(chlorides, sulphates, alcohol)1.0.0 -3.79313    0.76787   -4.94 9.01e-07 ***
poly(chlorides, sulphates, alcohol)0.1.0  6.73404    0.75321    8.94 < 2e-16 ***
poly(chlorides, sulphates, alcohol)0.0.1 11.58245    0.72439   15.99 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6937 on 1115 degrees of freedom
Multiple R-squared:  0.2824,    Adjusted R-squared:  0.2805
F-statistic: 146.3 on 3 and 1115 DF,  p-value: < 2.2e-16
```

## Model result

```
[1] "Head of the residual :"  
      residual  
1  -0.22339468  
2  -0.43876381  
4   0.61998616  
7  -0.12570398  
8   1.65622824  
11 -0.08371956  
  
Validation MSE: 0.6537456  
  
R^2 adjusted 0.2804642  
  
***** Prediction the test set *****  
Test MSE: 0.3938758  
> |
```

Polynomial regression has better MSE and slightly better  $R^2$  adjusted than simple linear regression. The error on the test set performs much better, therefore the model is not overfitting.

## Multiple Linear Regression

```
1 library(caTools)
2 library(ggplot2)
3 #Regression problem
4
5 #Upload the data - Heart disease UCI
6 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\winequality-red.csv")
7
8 #Check N/A value
9 print(any(is.na(data)))
10
11 #Check structure of the data
12 print(str(data))
13
14 #Train validation test split
15 set.seed(42)
16 spec <- c(train = .7, validate = .15, test = .15)
17
18 g <- sample(cut(
19   seq(nrow(data)),
20   nrow(data)*cumsum(c(0,spec)),
21   labels = names(spec)
22 ))
```

```
R - Global Environment
Data
• data      1599 obs. of 12 variables
• multi.lm_mod... List of 12
• res       List of 3
• residual  1119 obs. of 1 variable
• val.results 240 obs. of 2 variables
Values
g          Factor w/ 3 levels "train","validate...
spec       Named num [1:3] 0.7 0.15 0.15
test_mse   0.35148323161893
test_predict... Named num [1:240] 5.12 5.34 5.21 5.8...
val_mse    0.682382418518693
val.predicti... Named num [1:240] 5.24 5.05 5.08 5.3...
```

```

24 res <- split(data, g)
25
26 cat('\n Shape of train: \n', dim(res$train), '\n Shape of validation: \n', dim(res$validate), '\n Shape of test: \n', dim(res$test))
27
28 # Multiple Linear Regression
29 #Fit the model
30 multi.lm_model <- lm(quality ~ ., data = res$train)
31 print(summary(multi.lm_model))
32
33 #Plot residual check normality
34 #Grab residuals
35 residual <- residuals(multi.lm_model)
36
37 #Convert to dataframe and plot it
38 residual <- as.data.frame(residual)
39 print('Head of the residual :')
40 print(head(residual))
41
42 #ggplot
43 ggplot(residual, aes(residual)) + geom_histogram(fill = 'blue', alpha = 0.5)
44
45 #Predict on the validation set
46 val.prediction <- predict(multi.lm_model, res$validate)
47
48 #Result
49 val.results <- cbind(val.prediction, res$test$quality)
50 colnames(val.results) <- c('pred', 'valid')
51 val.results <- as.data.frame(val.results)
52
53 #Error
54 val_mse <- mean((val.results$valid - val.results$pred)^2)
55 cat('\nValidation MSE: ', val_mse, '\n')
56
57 #R^2 adjusted
58 cat('\nR^2 adjusted ', summary(multi.lm_model)$adj.r.squared, '\n')
59
60 #Predict on the test set
61 cat('\n***** Prediction the test set *****')
62 test_predictions <- predict(multi.lm_model, res$test)
63
64 #Error
65 test_mse <- mean((test_predictions - res$test$quality)^2)
66 cat('\n Test MSE: ', test_mse, '\n')
67

```

## Code output

```

> source("~/active-rstudio-document")
[1] FALSE
'data.frame': 1599 obs. of 12 variables:
 $ fixed.acidity : num 7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
 $ volatile.acidity : num 0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
 $ citric.acid : num 0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
 $ residual.sugar : num 1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 6.1 ...
 $ chlorides : num 0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
 $ free.sulfur.dioxide : num 11 25 15 17 11 13 15 15 9 17 ...
 $ total.sulfur.dioxide: num 34 67 54 60 34 40 59 21 18 102 ...
 $ density : num 0.998 0.997 0.997 0.998 0.998 ...
 $ pH : num 3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
 $ sulphates : num 0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57 0.8 ...
 $ alcohol : num 9.4 9.8 9.8 9.8 9.4 9.4 9.4 10 9.5 10.5 ...
 $ quality : int 5 5 5 6 5 5 5 7 7 5 ...
NULL

Shape of train:
1119 12
Shape of validation:
240 12
Shape of test:
240 12

```

```
Call:
lm(formula = quality ~ ., data = res$train)

Residuals:
    Min       1Q   Median       3Q      Max
-2.61208 -0.36766 -0.06288  0.46754  1.94936

Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  3.264e+01  2.536e+01   1.287   0.1982
fixed.acidity  4.199e-02  3.156e-02   1.331   0.1836
volatile.acidity -1.072e+00  1.464e-01  -7.321 4.73e-13 ***
citric.acid    -2.935e-01  1.777e-01  -1.652   0.0988 .
residual.sugar  3.115e-02  1.806e-02   1.725   0.0847 .
chlorides     -2.024e+00  5.078e-01  -3.985 7.19e-05 ***
free.sulfur.dioxide 1.293e-03  2.674e-03   0.484   0.6287
total.sulfur.dioxide -2.413e-03  9.154e-04  -2.636   0.0085 **
density       -2.886e+01  2.591e+01  -1.114   0.2656
pH            -3.794e-01  2.335e-01  -1.625   0.1044
sulphates     8.699e-01  1.337e-01   6.505 1.18e-10 ***
alcohol       2.819e-01  3.142e-02   8.974 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6616 on 1107 degrees of freedom
Multiple R-squared:  0.352,    Adjusted R-squared:  0.3456
F-statistic: 54.67 on 11 and 1107 DF,  p-value: < 2.2e-16
```

## Error

```
[1] "Head of the residual :"
```

	residual
1	-0.05065719
2	-0.15393853
4	0.29984570
7	-0.14394985
8	1.63422167
11	-0.07943643

```

Validation MSE: 0.6823824

R^2 adjusted 0.3455705

***** Prediction the test set *****
Test MSE: 0.3514832
```

Multiple Linear Regression is performed much better than simple linear regression and polynomial regression. Under both MSE and R\_squared adjusted. Again, the validation error is higher than the test error indicating the model is not overfitting.

## Natural Cubic Spline

```

1 library(caTools)
2 library(ggplot2)
3 library(splines)
4 #Regression problem
5
6 #Upload the data - Heart disease UCI
7 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\winequality-red.csv")
8
9 #Check N/A value
10 print(any(is.na(data)))
11
12 #Check structure of the data
13 print(str(data))
14
15 #Train validation test split
16 set.seed(42)
17 spec <- c(train = .7, validate = .15, test = .15)
18
19 g <- sample(cut(
20   seq(nrow(data)),
21   nrow(data)^cumsum(c(0, spec)),
22   labels = names(spec)
23 ))
```

Data	
data	1599 obs. of 12 variables
na_cub_splin...	List of 12
res	List of 3
residual	1119 obs. of 1 variable
val.results	240 obs. of 2 variables
Values	
g	Factor w/ 3 levels "train","validate...
spec	Named num [1:3] 0.7 0.15 0.15
test_mse	0.367998585750221
test_predict...	Named num [1:240] 5.38 5.1 5.28 6 5...
val_mse	0.686135348818267
val.predicti...	Named num [1:240] 5.52 5.2 5.2 5.25 ...



```

25 res <- split(data, g)
26
27 cat('\n Shape of train: \n', dim(res$train), '\n Shape of validation: \n', dim(res$validate), '\n Shape of test: \n', dim(res$test))
28
29 # Multiple Linear Regression
30 # Fit the model
31 na_cub_spline.lm_model <- lm(quality ~ ns(chlorides, df = 4) + ns(sulphates, df = 4) + ns(alc0hol, df = 4), data = res$train)
32 print(summary(na_cub_spline.lm_model))
33
34 # Plot residual check normality
35 # Grab residuals
36 residual <- residuals(na_cub_spline.lm_model)
37
38 # Convert to dataframe and plot it
39 residual <- as.data.frame(residual)
40 print('Head of the residual :')
41 print(head(residual))
42
43 # ggplot
44 ggplot(residual, aes(residual)) + geom_histogram(fill = 'blue', alpha = 0.5)
45

```

```

46 # Predict on the validation set
47 val.prediction <- predict(na_cub_spline.lm_model, res$validate)
48
49 # Result
50 val.results <- cbind(val.prediction, res$test$quality)
51 colnames(val.results) <- c('pred', 'valid')
52 val.results <- as.data.frame(val.results)
53
54 # Error
55 val_mse <- mean((val.results$valid - val.results$pred)^2)
56 cat('\n Validation MSE: ', val_mse, '\n')
57
58 # R^2 adjusted
59 cat('\n R^2 adjusted ', summary(na_cub_spline.lm_model)$adj.r.squared, '\n')
60
61 # Predict on the test set
62 cat('\n ***** Prediction the test set *****\n')
63 test_predictions <- predict(na_cub_spline.lm_model, res$test)
64
65 # Error
66 test_mse <- mean((test_predictions - res$test$quality)^2)
67 cat('\n Test MSE: ', test_mse, '\n')
68

```

## Code output

```

R 4.2.1 ~ /
> source("~/active-rstudio-document")
[1] FALSE
'data.frame': 1599 obs. of 12 variables:
 $ fixed.acidity      : num  7.4 7.8 7.8 11.2 7.4 7.4 7.9 7.3 7.8 7.5 ...
 $ volatile.acidity   : num  0.7 0.88 0.76 0.28 0.7 0.66 0.6 0.65 0.58 0.5 ...
 $ citric.acid        : num  0 0 0.04 0.56 0 0 0.06 0 0.02 0.36 ...
 $ residual.sugar     : num  1.9 2.6 2.3 1.9 1.9 1.8 1.6 1.2 2 6.1 ...
 $ chlorides          : num  0.076 0.098 0.092 0.075 0.076 0.075 0.069 0.065 0.073 0.071 ...
 $ free.sulfur.dioxide: num  11 25 15 17 11 13 15 15 9 17 ...
 $ total.sulfur.dioxide: num  34 67 54 60 34 40 59 21 18 102 ...
 $ density            : num  0.998 0.997 0.997 0.998 0.998 ...
 $ pH                 : num  3.51 3.2 3.26 3.16 3.51 3.51 3.3 3.39 3.36 3.35 ...
 $ sulphates          : num  0.56 0.68 0.65 0.58 0.56 0.56 0.46 0.47 0.57 0.8 ...
 $ alc0hol            : num  9.4 9.8 9.8 9.8 9.4 9.4 9.4 10 9.5 10.5 ...
 $ quality            : int   5 5 5 6 5 5 5 7 7 5 ...
NULL

Shape of train:
1119 12
Shape of validation:
240 12
Shape of test:
240 12

```

## Model summary

```
Call:
lm(formula = quality ~ ns(chlorides, df = 4) + ns(sulphates,
  df = 4) + ns(alcohol, df = 4), data = res$train)

Residuals:
    Min       1Q   Median       3Q      Max
-2.59935 -0.38198 -0.05765  0.45440  2.24151

Coefficients:
              Estimate Std. Error t value Pr(>|t|)
(Intercept)      4.8393     0.4523  10.699 < 2e-16 ***
ns(chlorides, df = 4)1 -0.2106     0.3544  -0.594  0.55251
ns(chlorides, df = 4)2 -0.8250     0.3052  -2.703  0.00698 **
ns(chlorides, df = 4)3 -1.0879     0.7755  -1.403  0.16094
ns(chlorides, df = 4)4 -0.6184     0.4443  -1.392  0.16425
ns(sulphates, df = 4)1  0.7166     0.1644   4.359 1.43e-05 ***
ns(sulphates, df = 4)2  1.2394     0.1784   6.946 6.38e-12 ***
ns(sulphates, df = 4)3  1.0995     0.4091   2.688  0.00730 **
ns(sulphates, df = 4)4  0.3081     0.3365   0.915  0.36021
ns(alcohol, df = 4)1    0.3184     0.2497   1.275  0.20251
ns(alcohol, df = 4)2    0.9366     0.1892   4.949 8.61e-07 ***
ns(alcohol, df = 4)3    1.2790     0.5641   2.267  0.02357 *
ns(alcohol, df = 4)4    1.2954     0.1978   6.549 8.84e-11 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.6769 on 1106 degrees of freedom
Multiple R-squared:  0.3224,    Adjusted R-squared:  0.315
F-statistic: 43.85 on 12 and 1106 DF,  p-value: < 2.2e-16
```

## Error

```
[1] "Head of the residual :"
```

	residual
1	-0.19909215
2	-0.55285117
4	0.64690796
7	0.08134628
8	1.91888783
11	-0.04543699

```
Validation MSE:  0.6861353

R^2 adjusted  0.3150192

***** Prediction the test set *****
Test MSE:  0.3679986
> |
```

Natural Cubic Spline with 3 predictor variables (chlorides, sulphates, and alcohol) along with 4 knots (df = 4) with almost the same MSE and R\_squared adjusted as the Multiple Linear

Regression. This indicates that the Natural Cubic Spline performs better than Multiple Linear Regression in terms of fewer independent variables. Hence lower computation expense compared to Multiple Linear Regression

## 2. Feature Selection/ Model Optimization

Breast Cancer dataset

### I. Forward/Backward Stepwise Selection

#### a. Forward Stepwise Selection

```

1 #Feature Selection/ Model optimization
2
3 library(leaps)
4 #Load the data
5 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\data.csv")
6
7 #Check what column have N/A value
8 colSums(is.na(data)) #Print this to check
9
10 #Remove N/A
11 data <- subset(data, select = -c(X))
12
13 #Subsets selection
14 regfit.fwd <- regsubsets(radius_mean ~ ., data = data, nvmax = 11, method = 'forward')
15 print(summary(regfit.fwd))
  
```

Global Environment

Variable	Value
data	569 obs. of 32 variables
na_cub_spli...	List of 12
regfit.fwd	List of 28
res	List of 3
residual	1119 obs. of 1 variable
val.results	240 obs. of 2 variables

Values

Variable	Value
g	Factor w/ 3 levels "train","valida...
spec	Named num [1:3] 0.7 0.15 0.15
test_mse	0.367998585750221
test_predict...	Named num [1:240] 5.38 5.1 5.28 6 ...
val_mse	0.686135348818267
val.predict...	Named num [1:240] 5.52 5.2 5.2 5.2...

Code output

```

R 4.2.1 ~ /
> source("~/active-rstudio-document")
Subset selection object
Call: regsubsets.formula(radius_mean ~ ., data = data, nvmax = 11,
  method = "forward")
31 Variables (and intercept)
      Forced in Forced out
id                FALSE    FALSE
diagnosisM        FALSE    FALSE
texture_mean      FALSE    FALSE
perimeter_mean    FALSE    FALSE
area_mean         FALSE    FALSE
smoothness_mean   FALSE    FALSE
compactness_mean  FALSE    FALSE
concavity_mean    FALSE    FALSE
concave.points_mean FALSE    FALSE
symmetry_mean     FALSE    FALSE
fractal_dimension_mean FALSE    FALSE
radius_se         FALSE    FALSE
texture_se        FALSE    FALSE
perimeter_se      FALSE    FALSE
area_se           FALSE    FALSE
smoothness_se     FALSE    FALSE
compactness_se    FALSE    FALSE
concavity_se      FALSE    FALSE
concave.points_se FALSE    FALSE
symmetry_se       FALSE    FALSE
fractal_dimension_se FALSE    FALSE
radius_worst      FALSE    FALSE
texture_worst     FALSE    FALSE
perimeter_worst   FALSE    FALSE
area_worst        FALSE    FALSE
smoothness_worst  FALSE    FALSE
compactness_worst FALSE    FALSE
concavity_worst   FALSE    FALSE
concave.points_worst FALSE    FALSE
symmetry_worst    FALSE    FALSE
fractal_dimension_worst FALSE    FALSE
1 subsets of each size up to 11
Selection Algorithm: forward
  
```

		id	diagnosisM	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "
		concavity_mean	concave.points_mean	symmetry_mean	fractal_dimension_mean	radius_se	texture_se	
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "
		perimeter_se	area_se	smoothness_se	compactness_se	concavity_se	concave.points_se	symmetry_se
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "
		fractal_dimension_se	radius_worst	texture_worst	perimeter_worst	area_worst	smoothness_worst	
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "
		compactness_worst	concavity_worst	concave.points_worst	symmetry_worst	fractal_dimension_worst		
1	(1)	" "	" "	" "	" "	" "	" "	" "
2	(1)	" "	" "	" "	" "	" "	" "	" "
3	(1)	" "	" "	" "	" "	" "	" "	" "
4	(1)	" "	" "	" "	" "	" "	" "	" "
5	(1)	" "	" "	" "	" "	" "	" "	" "
6	(1)	" "	" "	" "	" "	" "	" "	" "
7	(1)	" "	" "	" "	" "	" "	" "	" "
8	(1)	" "	" "	" "	" "	" "	" "	" "
9	(1)	" "	" "	" "	" "	" "	" "	" "
10	(1)	" "	" "	" "	" "	" "	" "	" "
11	(1)	" "	" "	" "	" "	" "	" "	" "

## b. Backward Stepwise Selection

### Code

```

18 #Backward
19 regfit.bwd <- regsubsets(radius_mean ~ ., data = data, nvmax = 11, method = "backward")
20 summary(regfit.bwd)
21
20:20 (Top Level) - R Script
Console Terminal Background Jobs
> #Backward
> regfit.bwd <- regsubsets(radius_mean ~ ., data = data, nvmax = 11, method = "backward")
> summary(regfit.bwd)

```

## Code output

```

Console Terminal Background Jobs
R 4.2.1 - ~/
Subset selection object
Call: regsubsets.formula(radius_mean ~ ., data = data, nvmax = 11,
  method = "backward")
31 Variables (and intercept)
Forced in Forced out
id FALSE FALSE
diagnosisM FALSE FALSE
texture_mean FALSE FALSE
perimeter_mean FALSE FALSE
area_mean FALSE FALSE
smoothness_mean FALSE FALSE
compactness_mean FALSE FALSE
concavity_mean FALSE FALSE
concave.points_mean FALSE FALSE
symmetry_mean FALSE FALSE
fractal_dimension_mean FALSE FALSE
radius_se FALSE FALSE
texture_se FALSE FALSE
perimeter_se FALSE FALSE
area_se FALSE FALSE
smoothness_se FALSE FALSE
compactness_se FALSE FALSE
concavity_se FALSE FALSE
concave.points_se FALSE FALSE
symmetry_se FALSE FALSE
fractal_dimension_se FALSE FALSE
radius_worst FALSE FALSE
texture_worst FALSE FALSE
perimeter_worst FALSE FALSE
area_worst FALSE FALSE
smoothness_worst FALSE FALSE
compactness_worst FALSE FALSE
concavity_worst FALSE FALSE
concave.points_worst FALSE FALSE
symmetry_worst FALSE FALSE
fractal_dimension_worst FALSE FALSE
1 subsets of each size up to 11
Selection Algorithm: backward

```

```

R 4.2.1 - ~/
1 (1) id diagnosisM texture_mean perimeter_mean area_mean smoothness_mean compactness_mean
2 (1) " " " " " " " " " " " "
3 (1) " " " " " " " " " " " "
4 (1) " " " " " " " " " " " "
5 (1) " " " " " " " " " " " "
6 (1) " " " " " " " " " " " "
7 (1) " " " " " " " " " " " "
8 (1) " " " " " " " " " " " "
9 (1) " " " " " " " " " " " "
10 (1) " " " " " " " " " " " "
11 (1) " " " " " " " " " " " "
1 (1) concavity_mean concave.points_mean symmetry_mean fractal_dimension_mean radius_se texture_se
2 (1) " " " " " " " " " " " "
3 (1) " " " " " " " " " " " "
4 (1) " " " " " " " " " " " "
5 (1) " " " " " " " " " " " "
6 (1) " " " " " " " " " " " "
7 (1) " " " " " " " " " " " "
8 (1) " " " " " " " " " " " "
9 (1) " " " " " " " " " " " "
10 (1) " " " " " " " " " " " "
11 (1) " " " " " " " " " " " "
1 (1) perimeter_se area_se smoothness_se compactness_se concavity_se concave.points_se symmetry_se
2 (1) " " " " " " " " " " " "
3 (1) " " " " " " " " " " " "
4 (1) " " " " " " " " " " " "
5 (1) " " " " " " " " " " " "
6 (1) " " " " " " " " " " " "
7 (1) " " " " " " " " " " " "
8 (1) " " " " " " " " " " " "
9 (1) " " " " " " " " " " " "
10 (1) " " " " " " " " " " " "
11 (1) " " " " " " " " " " " "

```

```

1 (1) fractal_dimension_se radius_worst texture_worst perimeter_worst area_worst smoothness_worst
2 (1) " " " " " " " " " "
3 (1) " " " " " " " " " "
4 (1) " " " " " " " " " "
5 (1) " " " " " " " " " "
6 (1) " " " " " " " " " "
7 (1) " " " " " " " " " "
8 (1) " " " " " " " " " "
9 (1) " " " " " " " " " "
10 (1) " " " " " " " " " "
11 (1) " " " " " " " " " "

1 (1) compactness_worst concavity_worst concave.points_worst symmetry_worst fractal_dimension_worst
2 (1) " " " " " " " " " "
3 (1) " " " " " " " " " "
4 (1) " " " " " " " " " "
5 (1) " " " " " " " " " "
6 (1) " " " " " " " " " "
7 (1) " " " " " " " " " "
8 (1) " " " " " " " " " "
9 (1) " " " " " " " " " "
10 (1) " " " " " " " " " "
11 (1) " " " " " " " " " "
>

```

## II. Forward/Backward Features

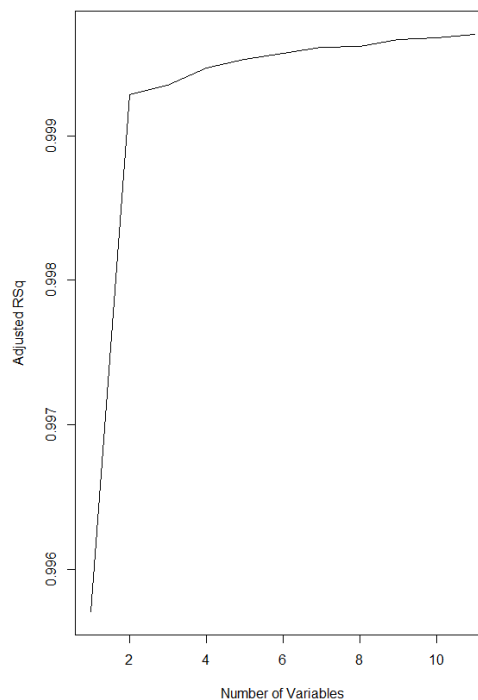
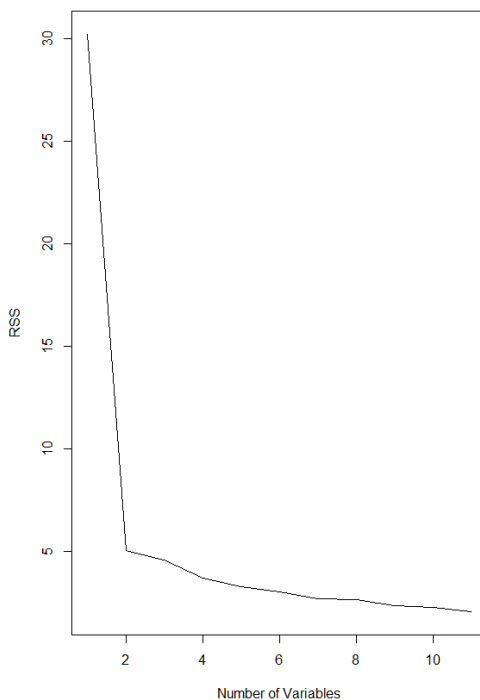
### a. Forward Features

```

22 #Part II
23 #Forward Features
24 reg.summaryfwd <- summary(regfit.fwd)
25 par(mfrow = c(1, 2))
26 plot(reg.summaryfwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
27 plot(reg.summaryfwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
28
29

```

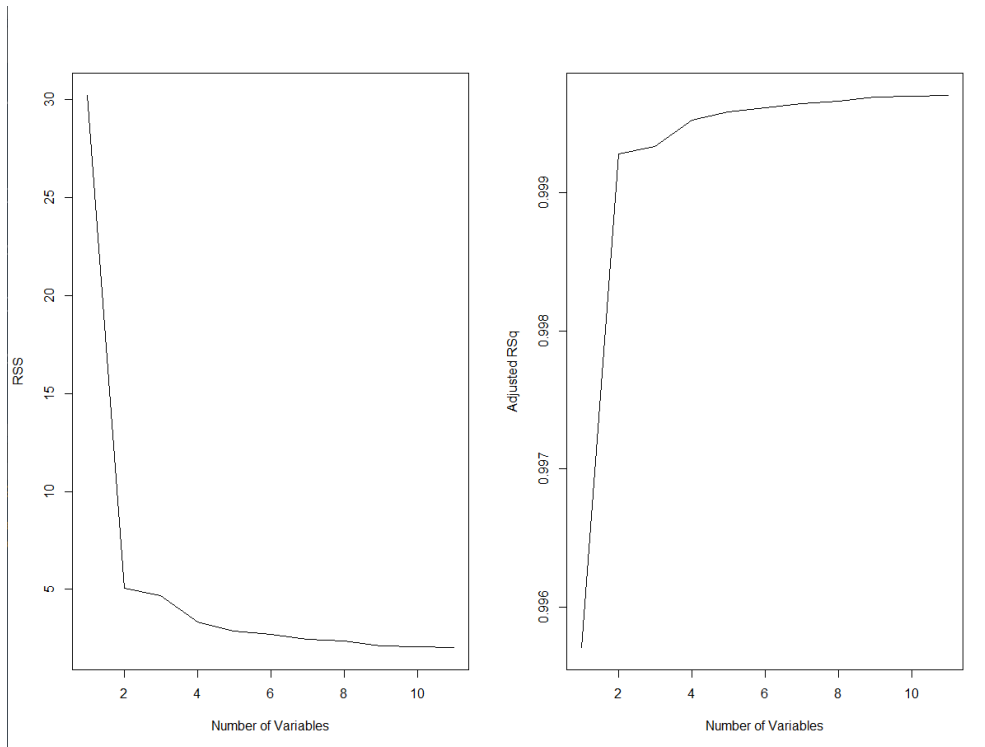
Code output



## b. Backward Features

```
29 #Backward Features
30 reg.summarybwd <- summary(regfit.bwd)
31 par(mfrow = c(1, 2))
32 plot(reg.summarybwd $rss, xlab = "Number of Variables", ylab = "RSS", type = "l")
33 plot(reg.summarybwd $adjr2, xlab = "Number of Variables", ylab = "Adjusted RSq", type = "l")
34
35
```

Code output



## III. PCA

Code

```
36 #Part III
37 #PCA and PLR
38
39 library(pls)
40 set.seed(42)
41
42 #Fit the model
43 pcr.fit <- pcr(radius_mean ~ ., data = data, scale = TRUE, validation = 'cv')
44 print(summary(pcr.fit))
45
46
```

Code output

```

> set.seed(42)
>
> #Fit the model
> pcr.fit <- pcr(radius_mean ~ ., data = data, scale = TRUE, validation = 'CV')
> print(summary(pcr.fit))
Data:   X dimension: 569 31
        Y dimension: 569 1
Fit method: svdpc
Number of components considered: 31
TRAINING: % variance explained

```

	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps
X	42.90	60.38	69.62	76.05	81.39	85.37	88.52	90.69	92.21
radius_mean	59.74	93.11	93.11	93.75	94.11	94.23	94.42	95.19	95.20

	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	16 comps	17 comps
X	93.44	94.55	95.53	96.47	97.31	97.99	98.46	98.75
radius_mean	97.34	98.00	98.04	98.30	98.41	98.44	98.49	98.51

	18 comps	19 comps	20 comps	21 comps	22 comps	23 comps	24 comps	25 comps
X	99.00	99.18	99.34	99.49	99.58	99.68	99.77	99.84
radius_mean	98.67	98.79	98.88	99.58	99.58	99.61	99.65	99.71

	26 comps	27 comps	28 comps	29 comps	30 comps	31 comps
X	99.90	99.95	99.97	99.99	100.00	100.00
radius_mean	99.86	99.87	99.90	99.95	99.97	99.97

```

NULL
>

```

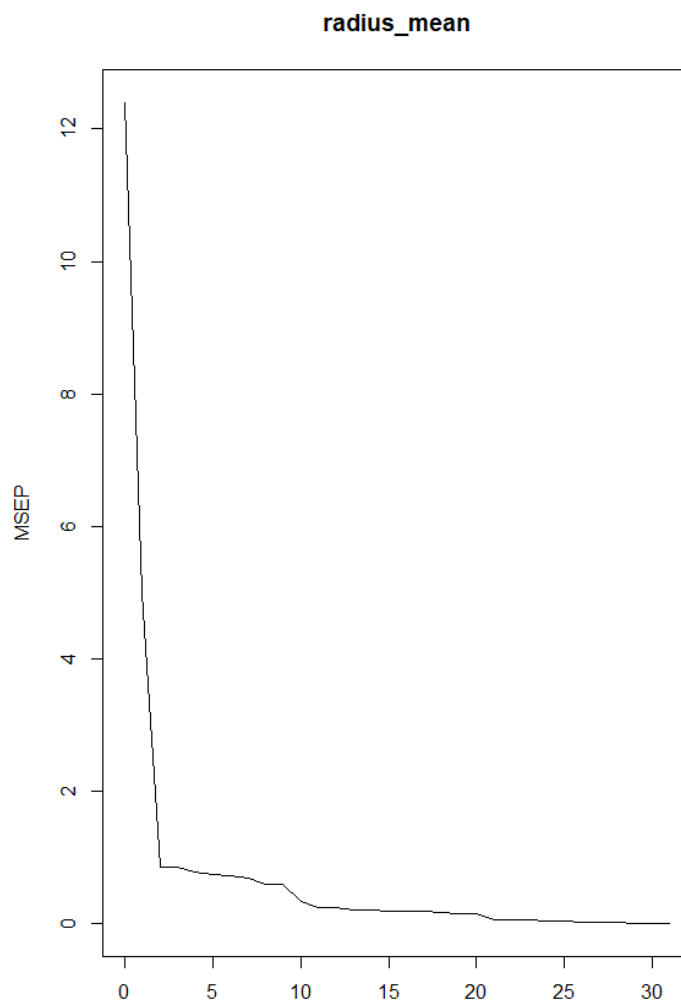
Code

```

R 4.2.1 ~ /
> validationplot(pcr.fit, val.type = "MSEP")
>

```

Code output

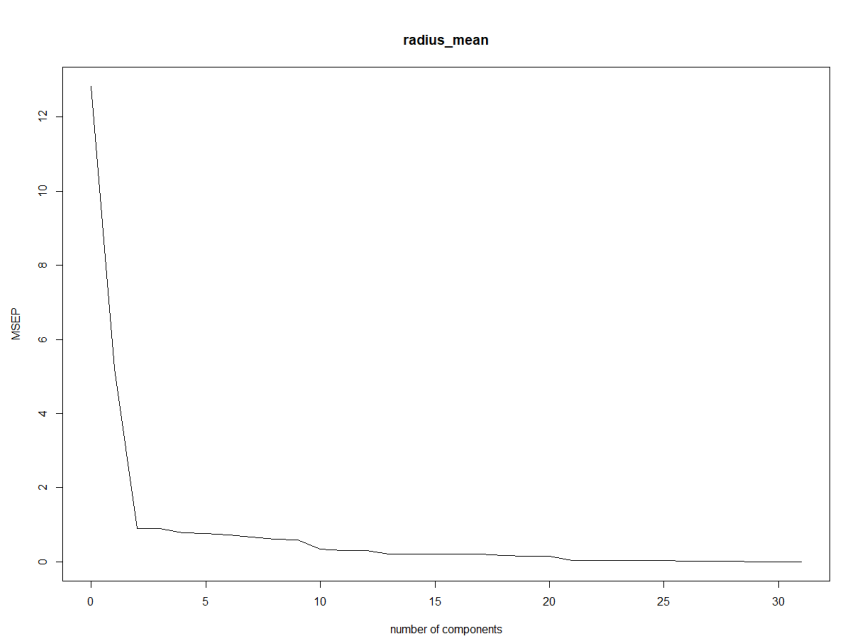




## Code

```
R 4.2.1 ~ /
> set.seed(42)
> spec <- c(train = .7, validate = .15, test = .15)
>
> g <- sample(cut(
+   seq(nrow(data)),
+   nrow(data)*cumsum(c(0,spec)),
+   labels = names(spec)
+ ))
>
> res <- split(data, g)
>
> pcr.fit <- pcr(radius_mean ~ ., data = res$train, scale = TRUE, validation = 'cv')
> validationplot(pcr.fit, val.type = "MSEP")
> |
```

## Code output



## Code output

```
R 4.2.1 ~ /
> #Predict the model on validation set
> pcr.pred <- predict(pcr.fit, res$validate, ncomp = 5)
> val_mse <- mean((pcr.pred - res$validate$radius_mean)^2)
> cat('\nvalidation MSE: ', val_mse, '\n')

Validation MSE: 0.6542805
>
>
> pcr.fit <- pcr(radius_mean ~ ., data = data, scale = TRUE, ncomp = 5)
> summary(pcr.fit)
Data:   X dimension: 569 31
       Y dimension: 569 1
Fit method: svdpc
Number of components considered: 5
TRAINING: % variance explained
      1 comps  2 comps  3 comps  4 comps  5 comps
X      42.90   60.38   69.62   76.05   81.39
radius_mean 59.74   93.11   93.11   93.75   94.11
> |
```

## Principle Component Analysis

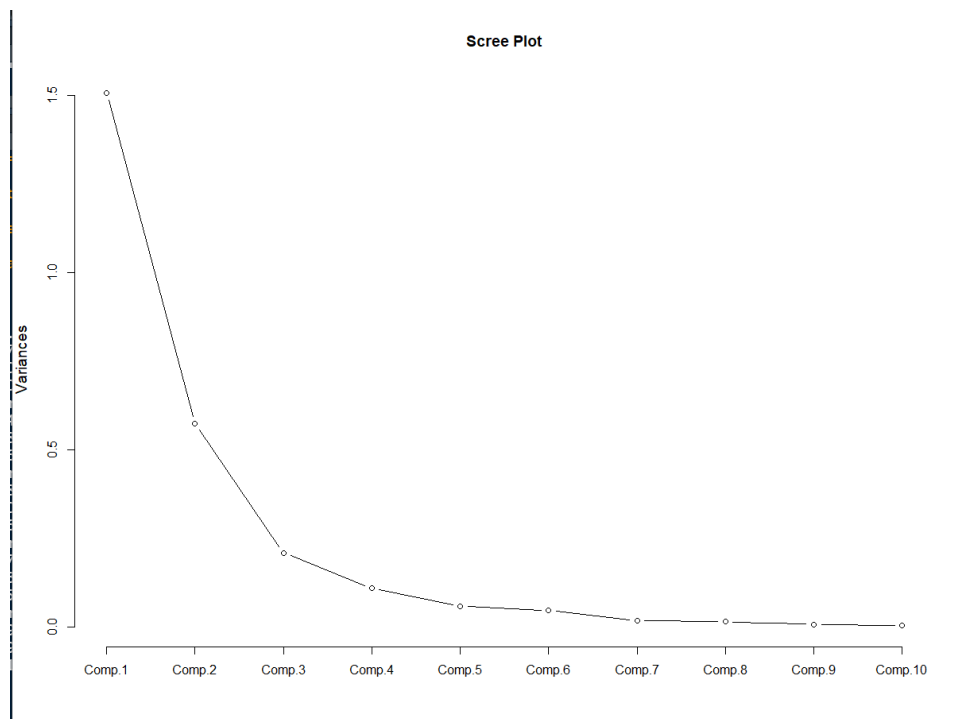
### Code

```
72
73 #PCA
74
75 data <- subset(data, select = -c(diagnosis))
76 #Standardize the variable
77 standardize_data <- scale(data)
78 #Correlation matrix
79 cor_matrix <- cor(standardize_data)
80 #PCA
81 pca_result <- princomp(cor_matrix)
82 #Print the summary
83 print(summary(pca_result))
84
85 #Scree plot
86 screeplot(pca_result, type = "lines", main = "Scree Plot", cex.axis = 1.2, cex.lab = 1.2)
87
88
```

### Code output

```
Importance of components:
      Comp.1      Comp.2      Comp.3      Comp.4      Comp.5      Comp.6      Comp.7
Standard deviation 1.2273776 0.7566901 0.45672005 0.32991045 0.24019324 0.2146628 0.128333718
Proportion of Variance 0.5903308 0.2243754 0.08174086 0.04265119 0.02260792 0.0180573 0.006453876
Cumulative Proportion 0.5903308 0.8147061 0.89644699 0.93909819 0.96170611 0.9797634 0.986217281
      Comp.8      Comp.9      Comp.10      Comp.11      Comp.12      Comp.13
Standard deviation 0.117594803 0.080254809 0.064600840 0.061917616 0.0487349582 0.0451179343
Proportion of Variance 0.005418952 0.002523953 0.001635368 0.001502337 0.0009307226 0.0007976961
Cumulative Proportion 0.991636233 0.994160186 0.995795553 0.997297891 0.9982286135 0.9990263096
      Comp.14      Comp.15      Comp.16      Comp.17      Comp.18      Comp.19
Standard deviation 0.0311585168 0.02816656 0.0163602157 1.123179e-02 9.575856e-03 8.830673e-03
Proportion of Variance 0.0003804456 0.00031089 0.0001048859 4.943528e-05 3.593306e-05 3.055812e-05
Cumulative Proportion 0.9994067553 0.99971765 0.9998225312 9.998720e-01 9.999079e-01 9.999385e-01
      Comp.20      Comp.21      Comp.22      Comp.23      Comp.24      Comp.25
Standard deviation 7.153067e-03 5.496602e-03 5.371526e-03 4.540499e-03 3.277884e-03 3.041872e-03
Proportion of Variance 2.005043e-05 1.183934e-05 1.130666e-05 8.078787e-06 4.210427e-06 3.625943e-06
Cumulative Proportion 9.999585e-01 9.999703e-01 9.999817e-01 9.999897e-01 9.999939e-01 9.999976e-01
      Comp.26      Comp.27      Comp.28      Comp.29      Comp.30      Comp.31
Standard deviation 2.002659e-03 1.433904e-03 3.198257e-04 1.646988e-04 8.327191e-05 0
Proportion of Variance 1.571639e-06 8.057105e-07 4.008352e-08 1.062967e-08 2.717290e-09 0
Cumulative Proportion 9.999991e-01 9.999999e-01 1.000000e+00 1.000000e+00 1.000000e+00 1
>
```

### Scree plot



Since 6 components capture 97% of the total variance. Hence select 6 PCs

```
> selected_components <- 1:6 # Adjust as needed
> loadings_selected <- loadings[, selected_components]
>
> # Print loadings
> print(loadings_selected)
```

	Comp.1	Comp.2	Comp.3	Comp.4	Comp.5	Comp.6
id	0.005275410	0.117520931	0.005999507	0.060862453	0.478220652	0.083898356
radius_mean	0.321496059	0.015241834	-0.022744524	0.007276042	0.015242645	0.033488737
texture_mean	0.075372239	0.111033932	0.150551922	-0.467816522	-0.239292189	-0.155096433
perimeter_mean	0.315065625	-0.003650746	-0.027636127	0.005202839	0.009906413	0.028365363
area_mean	0.316742118	0.028972666	-0.053412594	0.028599052	0.010671753	0.008782555
smoothness_mean	-0.024742995	-0.230758795	0.047659507	0.320697166	-0.131349892	-0.370426203
compactness_mean	0.076005866	-0.285758190	-0.028702588	-0.011870725	-0.048263909	-0.046502986
concavity_mean	0.148439860	-0.212642698	-0.094927033	-0.069337569	-0.004896528	-0.022104457
concave.points_mean	0.215277925	-0.151474340	-0.046346093	0.050816915	-0.055950582	-0.084578600
symmetry_mean	-0.037759041	-0.203722614	0.008908493	0.196954866	-0.347775237	0.225410820
fractal_dimension_mean	-0.205792130	-0.296678472	-0.056576036	0.038116744	-0.014618619	-0.142180719
radius_se	0.196164374	0.047007678	-0.269190878	0.109703856	-0.139080611	-0.095773645
texture_se	-0.107218897	0.140076418	-0.188151568	-0.250183339	-0.391001759	-0.206745374
perimeter_se	0.190771209	0.027076418	-0.277094199	0.079714690	-0.140272707	-0.069122772
area_se	0.231145441	0.061247143	-0.220693603	0.118499754	-0.071806532	-0.084418803
smoothness_se	-0.170772144	0.002608844	-0.258739282	0.125540095	-0.041843311	-0.385650835
compactness_se	-0.042412803	-0.235604257	-0.253941051	-0.236412365	0.068783897	0.093466835
concavity_se	-0.027237703	-0.206811404	-0.288120496	-0.253244576	0.140509773	0.109964924
concave.points_se	0.034752937	-0.157934388	-0.324610999	-0.115583928	0.059090685	-0.020590407
symmetry_se	-0.131410601	-0.020497029	-0.252145154	0.116919287	-0.399842307	0.359923346
fractal_dimension_se	-0.129889458	-0.203511611	-0.299235657	-0.187963264	0.134642008	-0.001714700
radius_worst	0.318051743	-0.002408842	0.025847354	0.015915974	0.001231180	0.005055467
texture_worst	0.071619856	0.073167730	0.261454523	-0.452216642	-0.256835218	-0.179836565
perimeter_worst	0.311376963	-0.024130719	0.018644195	0.004413788	-0.002219419	0.011130842
area_worst	0.311258663	0.015550245	-0.001918674	0.034893550	0.004243826	-0.020892397
smoothness_worst	-0.021763199	-0.238278407	0.241067655	0.222802201	-0.029879936	-0.413403455
compactness_worst	0.071405654	-0.309037963	0.146259435	-0.139611665	0.033960468	0.055266345
concavity_worst	0.114812239	-0.278328383	0.074223352	-0.171883896	0.076134438	0.052182604
concave.points_worst	0.202022267	-0.213855067	0.084136312	-0.022263882	0.023356840	-0.029454089
symmetry_worst	-0.001850243	-0.224726359	0.256603325	0.128834924	-0.286304187	0.415024460
fractal_dimension_worst	-0.079062657	-0.341519242	0.148406365	-0.104861644	0.092614311	-0.052821919

Each number in the PCs indicates the correlation between variables. For example, PC1 indicates slight correlation between perimeter\_mean (0.315), area\_mean (0.316), etc.

### 3. Classification

#### I.

#### Logistics Regression

Code

```
1 #Classification
2
3 #Upload data
4 data <- read.csv("C:\\Users\\hoang\\OneDrive\\Desktop\\data.csv")
5
6 #Drop N/A column
7 data <- subset(data, select = -c(X))
8
9 #Convert character data to numerical
10 data$diagnosis <- as.numeric(data$diagnosis == "M") # 'M' is mapped to 1 and 'B' is mapped to 0
11
12 #Train val test split
13 spec <- c(train = .7, validate = .15, test = .15)
14
15 g <- sample(cut(
16   seq(nrow(data)),
17   nrow(data)*cumsum(c(0,spec))),
18   labels = names(spec)
19 )
20
21 res <- split(data, g)
22
23 #I - Logisitcs Regression and LDA
24
25 #Logistics
26 #Fit the model
27 set.seed(42)
28
29 log.model <- glm(diagnosis ~ ., data = res$train, family = binomial(link = 'logit'))
30
```

```

31 print(summary(log.model))
32
33 #Predict on validation set
34 fitted.probabilities <- predict(log.model, newdata = res$validate, type = 'response' )
35 fitted.result <- ifelse(fitted.probabilities > 0.5, 1, 0)
36
37 #Validation error
38 misClassificError <- mean(fitted.result != res$validate$diagnosis)
39 cat('\nAccuracy of validation set: ', 1 - misClassificError, '\n')
40
41 #Missclassification table
42 table(res$validate$diagnosis, fitted.probabilities > 0.5)
43
44 #Predict on test set
45 fitted.probabilities <- predict(log.model, newdata = res$test, type = 'response' )
46 fitted.result <- ifelse(fitted.probabilities > 0.5, 1, 0)
47
48 #Validation error
49 misClassificError <- mean(fitted.result != res$test$diagnosis)
50 cat('\nAccuracy of test set: ', 1 - misClassificError, '\n')
51
52

```

Code output:

```

Call:
glm(formula = diagnosis ~ ., family = binomial(link = "logit"),
    data = res$train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-1.334e-04 -2.100e-08 -2.100e-08  2.100e-08  1.200e-04

Coefficients:
            Estimate Std. Error z value Pr(>|z|)
(Intercept) -1.009e+03  1.583e+06  -0.001  0.999
id           -4.470e-08  5.937e-04   0.000  1.000
radius_mean  -4.186e+01  6.050e+05   0.000  1.000
texture_mean  -1.685e+00  2.883e+04   0.000  1.000
perimeter_mean 5.593e+00  9.852e+04   0.000  1.000
area_mean     -8.652e-02  2.226e+03   0.000  1.000
smoothness_mean 3.082e+03  1.064e+07   0.000  1.000
compactness_mean -1.228e+03  3.274e+06   0.000  1.000
concavity_mean  1.855e+03  3.326e+06   0.001  1.000
concave.points_mean -1.372e+03  6.192e+06   0.000  1.000
symmetry_mean  -3.406e+02  1.159e+06   0.000  1.000
fractal_dimension_mean 1.491e+03  1.276e+07   0.000  1.000
radius_se     -1.840e+02  1.758e+06   0.000  1.000
texture_se     -8.693e+01  1.333e+05  -0.001  0.999
perimeter_se   8.483e+00  1.547e+05   0.000  1.000
area_se        4.594e+00  8.700e+03   0.001  1.000
smoothness_se  1.944e+04  2.180e+07   0.001  0.999
compactness_se 6.588e+03  8.051e+06   0.001  0.999
concavity_se   -3.773e+03  3.410e+06  -0.001  0.999
concave.points_se 2.823e+03  1.585e+07   0.000  1.000
symmetry_se    -3.038e+03  3.536e+06  -0.001  0.999
fractal_dimension_se -6.164e+04  6.191e+07  -0.001  0.999
radius_worst   6.326e+01  1.575e+05   0.000  1.000
texture_worst   9.855e+00  2.507e+04   0.000  1.000

perimeter_worst -5.479e+00  1.210e+04   0.000  1.000
area_worst      -1.329e-01  1.161e+03   0.000  1.000
smoothness_worst -2.071e+03  2.701e+06  -0.001  0.999
compactness_worst -6.782e+02  1.072e+06  -0.001  0.999
concavity_worst  1.743e+02  8.743e+05   0.000  1.000
concave.points_worst 1.454e+03  1.954e+06   0.001  0.999
symmetry_worst   4.166e+02  8.799e+05   0.000  1.000
fractal_dimension_worst 5.494e+03  5.694e+06   0.001  0.999

(Dispersion parameter for binomial family taken to be 1)

    Null deviance: 5.2836e+02  on 397  degrees of freedom
Residual deviance: 1.2192e-07  on 366  degrees of freedom
AIC: 64

Number of Fisher Scoring iterations: 25

>
> #Predict on validation set
> fitted.probabilities <- predict(log.model, newdata = res$validate, type = 'response' )
> fitted.result <- ifelse(fitted.probabilities > 0.5, 1, 0)
>
> #Validation error
> misClassificError <- mean(fitted.result != res$validate$diagnosis)
> cat('\nAccuracy of validation set: ', 1 - misClassificError, '\n')

Accuracy of validation set:  0.9411765

>
> #Missclassification table
> table(res$validate$diagnosis, fitted.probabilities > 0.5)

    FALSE TRUE
0       51   1
1        4  29
>

```

```

> #Predict on test set
> fitted.probabilities <- predict(log.model, newdata = res$test, type = 'response' )
> fitted.result <- ifelse(fitted.probabilities > 0.5, 1, 0)
>
> #Validation error
> misClassificError <- mean(fitted.result != res$test$diagnosis)
> cat('\nAccuracy of test set: ', 1 - misClassificError, '\n')

Accuracy of test set:  0.9534884
> |

```

The accuracy of the test set is higher than validation set indicate the model is not overfitting.

## Linear Discriminant Analysis

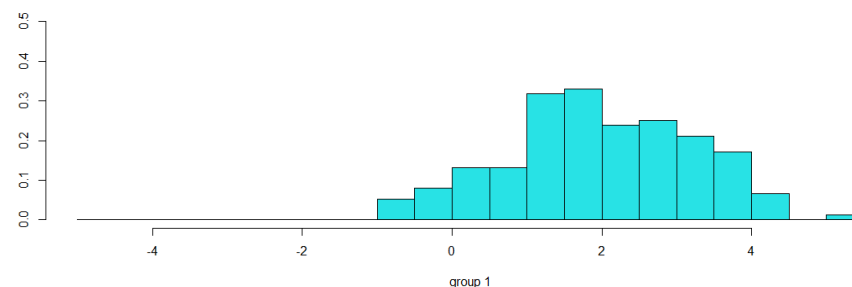
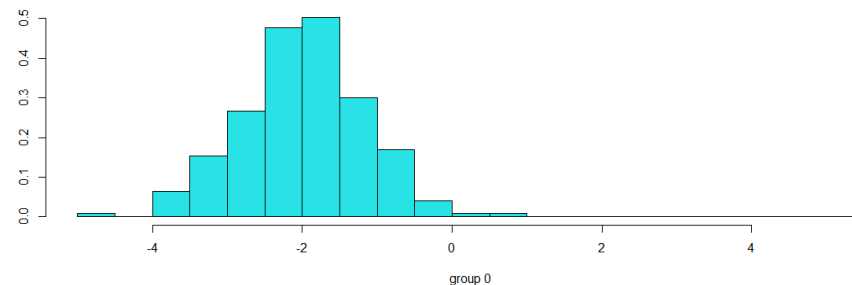
### Code

```

53 #LDA
54 library(MASS)
55
56 set.seed(42)
57 #Fit the model
58 lda.model <- lda(diagnosis ~., data = res$train)
59 print(lda.model)
60
61 #Plot
62 plot(lda.model)
63 |
64 #Predict on validation set
65 lda.pred <- predict(lda.model, res$validate)
66 # Extract predicted classes
67 predicted_classes <- lda.pred$class
68
69 # Display the confusion matrix (if you have actual classes for comparison)
70 confusion_matrix <- table(res$validate$diagnosis, predicted_classes)
71 cat("Confusion Matrix:\n", confusion_matrix, "\n")
72
73 # Calculate classification accuracy
74 accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
75 cat("Classification Accuracy: ", accuracy, "\n")
76
77
78 #Predict on test set
79 lda.pred <- predict(lda.model, res$test)
80 # Extract predicted classes
81 predicted_classes <- lda.pred$class
82
83 # Display the confusion matrix (if you have actual classes for comparison)
84 confusion_matrix <- table(res$test$diagnosis, predicted_classes)
85 cat("Confusion Matrix:\n", confusion_matrix, "\n")
86
87 # Calculate classification accuracy
88 accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
89 cat("Classification Accuracy: ", accuracy, "\n")

```

### Plot



## Code output

```
R421 ~/?
> library(MASS)
>
> set.seed(42)
> #Fit the model
> lda.model <- lda(diagnosis ~., data = res$train)
> print(lda.model)
call:
lda(diagnosis ~ ., data = res$train)

Prior probabilities of groups:
      0      1
0.620603 0.379397

Group means:
      id radius_mean texture_mean perimeter_mean area_mean smoothness_mean compactness_mean
0 27654428    12.24777    17.92721     78.68644   470.6964    0.09206599    0.0790534
1 45803457    17.43517    21.65278    115.32530   976.2066    0.10369987    0.1489775
      concavity_mean concave.points_mean symmetry_mean fractal_dimension_mean radius_se texture_se perimeter_se
0    0.0441974    0.02575779    0.1758611    0.06246134 0.2828069    1.206489    1.964308
1    0.1667068    0.09011570    0.1934166    0.06317550 0.6142821    1.191575    4.327192
      area_se smoothness_se compactness_se concavity_se concave.points_se symmetry_se fractal_dimension_se
0 21.17406    0.007100826    0.02029681    0.02399666    0.009588166 0.02029611    0.003453888
1 73.44172    0.006920543    0.03290591    0.04327523    0.015030437 0.02037019    0.004137848
      radius_worst texture_worst perimeter_worst area_worst smoothness_worst compactness_worst concavity_worst
0    13.49207    23.50802     87.58028    568.1206    0.1246253    0.1778482    0.1579797
1    21.17397    29.27179    141.75411    1433.1132    0.1466191    0.3816534    0.4644542
      concave.points_worst symmetry_worst fractal_dimension_worst
0    0.07347489    0.2713745    0.07864603
1    0.18469238    0.3227695    0.09249344
```

### Coefficients of linear discriminants:

```
LD1
id -2.225837e-10
radius_mean -1.002961e+00
texture_mean 4.101809e-02
perimeter_mean 1.600460e-01
area_mean -1.723805e-03
smoothness_mean -1.940282e-01
compactness_mean -2.751007e+01
concavity_mean 8.681797e+00
concave.points_mean 8.999125e+00
symmetry_mean -2.325073e+00
fractal_dimension_mean 1.985320e+01
radius_se 4.628067e+00
texture_se -6.454275e-02
perimeter_se -3.408389e-01
area_se -7.539064e-03
smoothness_se 6.789281e+01
compactness_se 1.144721e+01
concavity_se -1.805898e+01
concave.points_se 2.392216e+01
symmetry_se 1.512903e+01
fractal_dimension_se -5.511060e+01
radius_worst 6.297169e-01
texture_worst 1.330730e-02
perimeter_worst -6.326132e-03
area_worst -2.971028e-03
smoothness_worst 4.101605e+00
compactness_worst 1.063463e+00
concavity_worst 1.745605e+00
concave.points_worst 8.486796e+00
symmetry_worst 2.197523e+00
fractal_dimension_worst 1.695436e+01
>
```

```

> #Plot
> plot(lda.model)
>
> #Predict on validation set
> lda.pred <- predict(lda.model, res$validate)
> # Extract predicted classes
> predicted_classes <- lda.pred$class
>
> # Display the confusion matrix (if you have actual classes for comparison)
> confusion_matrix <- table(res$validate$diagnosis, predicted_classes)
> cat("Confusion Matrix:\n", confusion_matrix, "\n")
Confusion Matrix:
 52 6 0 27
>
> # Calculate classification accuracy
> accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
> cat("Classification Accuracy: ", accuracy, "\n")
Classification Accuracy: 0.9294118
>
>
> #Predict on test set
> lda.pred <- predict(lda.model, res$test)
> # Extract predicted classes
> predicted_classes <- lda.pred$class
>
> # Display the confusion matrix (if you have actual classes for comparison)
> confusion_matrix <- table(res$test$diagnosis, predicted_classes)
> cat("Confusion Matrix:\n", confusion_matrix, "\n")
Confusion Matrix:
 58 5 0 23
>
> # Calculate classification accuracy
> accuracy <- sum(diag(confusion_matrix)) / sum(confusion_matrix)
> cat("Classification Accuracy: ", accuracy, "\n")
Classification Accuracy: 0.9418605
>

```

Again, the accuracy of the test set is higher than the validation set indicating the LDA is not overfitting.

## II. Tree classifier

a.

Code

```

92 #II - Classification Trees
93 library(rpart)
94 set.seed(42)
95
96 #Fit the model
97 tree.model <- rpart(diagnosis ~ ., method = 'class', data = res$train)
98 print(summary(tree.model))
99

```

Code output

```

Call:
rpart(formula = diagnosis ~ ., data = res$train, method = "class")
n= 398

   CP nsplit rel error    xerror    xstd
1 0.81456954   0 1.0000000 1.0000000 0.06410892
2 0.07284768   1 0.1854305 0.2185430 0.03643216
3 0.01000000   2 0.1125828 0.1721854 0.03264673

Variable importance
concave.points_worst    perimeter_worst    radius_worst    concave.points_mean    concavity_mean
               19                15                15                15                14
      concavity_worst      area_worst      area_mean    perimeter_mean      radius_mean
               13                 3                 2                 2                 2

```



```

Node number 1: 398 observations, complexity param=0.8145695
predicted class=0 expected loss=0.379397 P(node) =1
class counts: 247 151
probabilities: 0.621 0.379
left son=2 (265 obs) right son=3 (133 obs)
Primary splits:
  concave.points_worst < 0.14655 to the left, improve=135.7905, (0 missing)
  concave.points_mean < 0.0501 to the left, improve=130.8744, (0 missing)
  area_worst < 884.55 to the left, improve=129.0354, (0 missing)
  perimeter_worst < 117.45 to the left, improve=128.8057, (0 missing)
  radius_worst < 16.795 to the left, improve=125.3750, (0 missing)
Surrogate splits:
  concave.points_mean < 0.059615 to the left, agree=0.935, adj=0.805, (0 split)
  concavity_mean < 0.1033 to the left, agree=0.927, adj=0.782, (0 split)
  perimeter_worst < 117.45 to the left, agree=0.907, adj=0.722, (0 split)
  concavity_worst < 0.344 to the left, agree=0.905, adj=0.714, (0 split)
  radius_worst < 17.29 to the left, agree=0.889, adj=0.669, (0 split)

Node number 2: 265 observations, complexity param=0.07284768
predicted class=0 expected loss=0.08679245 P(node) =0.6658291
class counts: 242 23
probabilities: 0.913 0.087
left son=4 (248 obs) right son=5 (17 obs)
Primary splits:
  area_worst < 929.8 to the left, improve=19.71960, (0 missing)
  radius_worst < 17.54 to the left, improve=17.93576, (0 missing)
  perimeter_worst < 116.05 to the left, improve=17.54312, (0 missing)
  area_se < 52.495 to the left, improve=17.33304, (0 missing)
  area_mean < 696.25 to the left, improve=16.10169, (0 missing)
Surrogate splits:
  radius_worst < 17.54 to the left, agree=0.996, adj=0.941, (0 split)
  area_mean < 794.25 to the left, agree=0.985, adj=0.765, (0 split)
  perimeter_worst < 115.45 to the left, agree=0.985, adj=0.765, (0 split)
  radius_mean < 15.92 to the left, agree=0.981, adj=0.706, (0 split)
  perimeter_mean < 103.55 to the left, agree=0.981, adj=0.706, (0 split)

Node number 3: 133 observations
predicted class=1 expected loss=0.03759398 P(node) =0.3341709
class counts: 5 128
probabilities: 0.038 0.962

Node number 4: 248 observations
predicted class=0 expected loss=0.03629032 P(node) =0.6231156
class counts: 239 9
probabilities: 0.964 0.036

Node number 5: 17 observations
predicted class=1 expected loss=0.1764706 P(node) =0.04271357
class counts: 3 14
probabilities: 0.176 0.824

n= 398

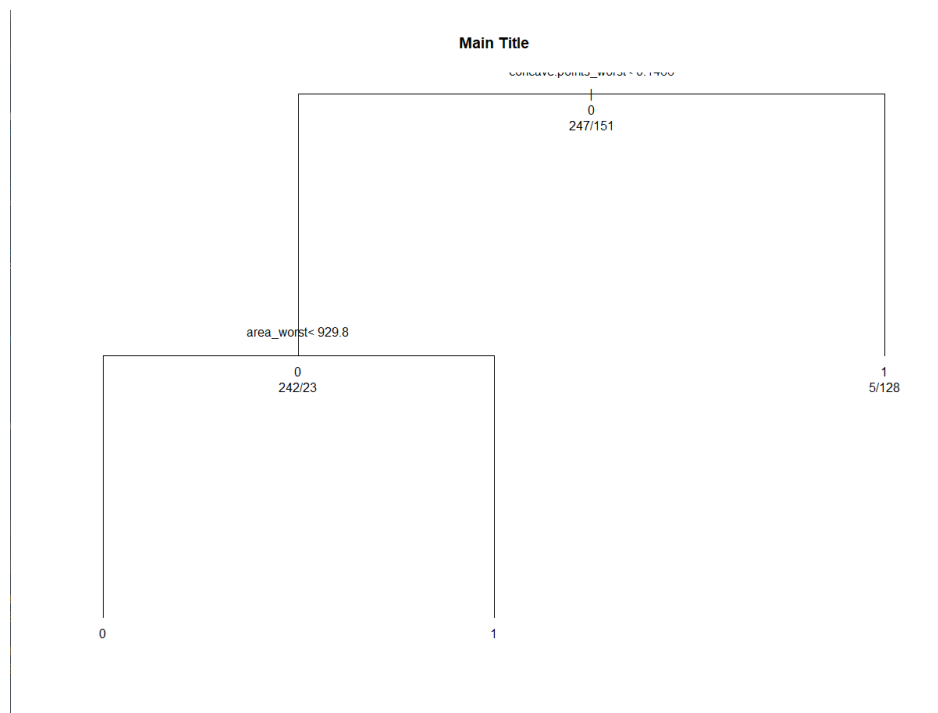
node), split, n, loss, yval, (yprob)
* denotes terminal node

1) root 398 151 0 (0.62060302 0.37939698)
2) concave.points_worst< 0.14655 265 23 0 (0.91320755 0.08679245)
4) area_worst< 929.8 248 9 0 (0.96370968 0.03629032) *
5) area_worst>=929.8 17 3 1 (0.17647059 0.82352941) *
3) concave.points_worst>=0.14655 133 5 1 (0.03759398 0.96240602) *
>

```

b.

Plot





### III. Support Vector Classifier

#### Code and output

```
R 4.2.1 - ~/
> #III - Support Vector Classifier
> library(e1071)
> set.seed(42)
>
> #Fit the model
> svm.model <- svm(diagnosis ~ ., data = data, kernel = 'linear', cost = 10, scale = TRUE)
>
> #Plot
> plot(svm.model, data)
>
> #Hyper parameter tuning
> tune.out <- tune(svm, diagnosis ~ ., data = data, kernel = "linear",
+               ranges = list(cost = c(0.001, 0.01, 0.1, 1, 5, 10, 100)))

WARNING: reaching max number of iterations
WARNING: reaching max number of iterations
WARNING: reaching max number of iterations
WARNING: reaching max number of iterations
WARNING: reaching max number of iterations
>
> print(summary(tune.out))

Parameter tuning of 'svm':

- sampling method: 10-fold cross validation

- best parameters:
  cost
  0.1

- best performance: 0.06389198

- Detailed performance results:
  cost      error dispersion
1 1e-03 0.07390136 0.01209621
2 1e-02 0.06650210 0.01225819
3 1e-01 0.06389198 0.01419311
4 1e+00 0.06537363 0.01475612
5 5e+00 0.06553679 0.01478780
6 1e+01 0.06590374 0.01472896
7 1e+02 0.06639767 0.01473196

>
> bestmod <- tune.out$best.model
> print(summary(bestmod))

Call:
best.tune(METHOD = svm, train.x = diagnosis ~ ., data = data, ranges = list(cost = c(0.001,
  0.01, 0.1, 1, 5, 10, 100)), kernel = "linear")

Parameters:
  SVM-Type:  eps-regression
  SVM-Kernel: linear
        cost: 0.1
        gamma: 0.03225806
        epsilon: 0.1

Number of Support Vectors: 455
```

#### 4. Models

- a. **Multiple Linear Regression.** Because multiple linear regression allows modeling the relationship between a continuous response variable (like profits) and multiple explanatory variables (like inventory, employees, budget etc). It can determine how strongly each factor correlates with profits (the measure of success). Also, multiple regression allows quantifying the correlation and predictive relationships between inventory, employees, budget, and other numerical factors to the level of business success measured by profits. This can give actionable and data-driven insights to my friend on what drives success in their industry. The interpretation is also straightforward.
- b. **Logistics Regression.** Because logistic regression is well-suited for binary outcomes, such as whether or not a shopper is likely to visit the store (1 for likely, 0 for not likely). In this case, the outcome variable can be binary based on the likelihood response in the survey data. Also, logistic regression provides interpretable coefficients, making it easier to understand the impact of predictor variables on the likelihood of shoppers visiting the store. This can be valuable when explaining the model to stakeholders.
- c. **Principal Component Analysis or Factor Analysis.** Because this scenario involves dimensionality reduction techniques, and PCA is specifically designed to reduce the dimensionality of the data while preserving as much of the original variance as possible. With thousands of features, PCA can transform the data into a lower-dimensional space, capturing the most important patterns and reducing computational complexity. PCA creates new features, called principal components, that are linear combinations of the original features. These components are orthogonal and ordered by their importance. The resulting principal components can provide insight into the most significant patterns in the data, aiding interpretability.
- d. **Random Forest.** First of all, because I'm a big fan of random forest when it comes to machine learning problem, random forest is the first thing that comes to my mind. Random Forest can handle complex relationships in the data and is robust to noisy and imprecise features. It builds multiple decision trees and combines their prediction. I'd choose Random Forest because of its simplicity and efficiency.